

# CIS 5500 Course Project Guidelines & Instructions

## TABLE OF CONTENTS

### A. Overview

- |                 |   |
|-----------------|---|
| 1. Description  | 2 |
| 2. Deliverables | 3 |

### B. Milestones and final presentation

- |   |    |
|---|----|
| 1. Milestone 0 – Team Members   | 5  |
| 2. Milestone 1 – Project Proposal   | 5  |
| 3. Milestone 2 – Project Outline  | 6  |
| 4. Milestone 3 – Database Population                                      | 8  |
| 5. Milestone 4 – Project Mentor Check-In and API specification submission | 9  |
| 6. Milestone 5 – Final Report, Demo Video, and Code                       | 10 |
| 7. Final Demo   | 11 |

### C. Grading

- |             |    |
|-------------|----|
| 1. Criteria | 12 |
| 2. Awards   | 13 |

### D. Resources

- |                      |    |
|----------------------|----|
| 1. Dataset Resources | 14 |
| 2. Example projects  | 14 |
| 3. Useful Tools      | 15 |

- |                      |    |
|----------------------|----|
| E. Plagiarism Policy | 17 |
|----------------------|----|

# OVERVIEW

## 1. Description

The goal of this project is to apply the concepts you've learned in this course to solve unstructured problems with limited hands-on guidance. In teams of four, you will first identify at least [two large](#), overlapping, and interesting datasets. You will import and integrate them into a [SQL relational database](#) and then design and build a web application that interacts with your database.

Some of the project requirements include:

- Wrangling and Cleaning data
  - Using programming language or excel to clean out anomalies in data (This will help in robustness of your web application and make it easier to debug code)
- Performing entity resolution
- Designing schema
  - Using Normal Forms learnt in class to reduce the tables into 3NF or BCNF
- Writing SQL queries
  - At least 2 complex queries (complex - use of CTEs, subqueries, multiple joins, aggregators - these are the queries you can optimize the most eventually)
  - Queries related to functionalities in your web app.
- Choosing database indexes
- Designing and Determining Web app functionality and pages
  - We require at least 3 pages with different functionalities
- Optimizing SQL queries
  - Using the concepts in class optimize queries written for pages to reduce latency in your web page.

If you use this as an opportunity to build something you're proud of, you are welcome to use the final product to demonstrate your skills to potential employers.

We also want you to enjoy completing the project. This is one of the most open ended projects in the Penn computer science curriculum. You can build just about any game, application, or website you want, as long as it's built on top of a SQL database which is queried using interesting queries and has multiple distinct pages. For example, past students have built trading card games, recipe recommendation platforms, and many other interesting apps. Check out the resources section for more inspiration.

In this document, you will find deliverables you'll produce while completing the project, each project milestone described in detail, the grading criteria and several

resources that might benefit you, including dataset resources, descriptions of past students' projects, and links to useful tools.

You will also be assigned a TA as a project mentor, who will be your primary point-of contact on staff for questions about your project. You may also post your questions relating to this project on Ed. Good luck and have fun!

## 2. Deliverables

The table below provides a brief description of each deliverable you'll need to submit to complete the project. Please see the milestone section for more detailed instructions.

<b>Deliverable Name</b>	<b>Description</b>	<b>Milestone</b>
<b>Project Group</b>	Form a team with diverse skills so everyone can contribute (Note: At the end of the project, we will take feedback from the team for each person's contribution and adjust the grades accordingly)	Milestone 0
<b>Project Proposal</b>	A 1-2 page document that identifies the datasets you will use, gives a rough idea of what your application may do, and demonstrates you have performed some basic descriptive or exploratory analysis on the data	Milestone 1
<b>Project Outline</b>	A 2-3 page document that describes what your application will do (pages, functionality), explains its significance, your motivation for the project, and provides the schema your database will implement	Milestone 2
<b>Database Population</b>	A 3NF- or BCNF-normalized MySQL database hosted on AWS RDS that contains all your data	Milestone 3 Part a
<b>SQL Queries</b>	A list of SQL queries that can run on your database with short explanations of what each query is supposed to do (with at least 2 complex queries)	Milestone 3 Part b
<b>Mentor Meeting</b>	An informal Zoom meeting with your project mentor where you demonstrate that you have implemented some basic functionality of your application	Milestone 4 Part a

<b>API Specs</b>	A small document which describes each route in your web application server; the description should include: the request path, request params, and response params.	Milestone 4 Part b
<b>Final Report</b>	A 6-10 page document that thoroughly describes the problem you tried to solve, your application functionality, your database design, your query optimization efforts, and more	Milestone 5 Part a
<b>Application Code</b>	A zip file that contains your application code, a list of dependencies, instructions for building the application and any code you wrote to populate the database or wrangle data	Milestone 5 Part b
<b>Application Demo</b>	A 2-4 minute screen-captured video that demonstrates your application's main functionalities, optimisation results and includes narration from at least one group member. (Upload on your UPenn account Google drive, make the link accessible to UPenn, and send us the link.)	Milestone 5 Part c
<b>Final Demo</b>	A video conference call with your project mentor and one other TA where you demonstrate functionality of your application. We use this to test edge cases and ensure your video demo wasn't heavily edited.	Final

# MILESTONES & FINAL PRESENTATION

## Milestone 0 – Project Group

Form a team of 3 or 4 (no more, no less).

The initial step is to select your teammates. It is important that you be able to work well together as a team so that everyone contributes to the project, and that the skillset of the group represent the project scope:

database cleaning and importing, creating a relational database, developing a web interface over a relational database, and optimizing queries of the database.

Factors that make it difficult to work together include dramatically different time zones/work schedules and intolerance of different work styles – you must learn to work together, and this will be reflected in a participation score in the final project grade.

We want to ensure that individual contributions are also graded.

## Milestone 1 – Project Proposal

Develop an initial idea.

Brainstorm ideas and applications, and search for large enough datasets to support the envisioned application. Remember your application must have at least 3 distinct pages and a SQL database with several tables. Your datasets must be:

1. **Large** -- As a rule of thumb, your final database should contain on the order of tens of thousands of instances or more (depending on the size of each instance). (Remember to use at least 2 large datasets to form your database; by large we mean more than 10,000 rows)
2. **Overlapping** -- Your application will need to include queries that require information from both datasets. This means the datasets need to be related and probably need to contain references to the same entities.

**Before finalizing your dataset choices, you should conduct some basic exploratory data analysis (EDA).** Poke around, compute summary statistics, and try to get a sense of how clean, large, and complete the data is. These factors will affect how difficult it will be to clean/pre-process the data, create and query a database, and perform entity resolution later. We provide an introduction to Pandas--a Python library for data analysis-- in [this tutorial](#) and demonstrate how to perform EDA with Pandas in [this tutorial](#).

Next, write your project proposal. The project proposal should contain the following:

1. A list of group members, email addresses, and GitHub usernames
2. A description of application/website idea
3. For each dataset you've chosen:
  - a. A 1-2 sentence description of the dataset
  - b. A link to where you found the dataset
  - c. If you're scraping the data, a description of how you will scrape it
  - d. If you're not scraping the data:
    - i. relevant size statistics (e.g. For a table, mb/gb, number of rows, and number of attributes. For a graph, mb, number of nodes, and number of edges)
    - ii. summary statistics of several attributes (e.g. report mean, standard deviation)
3. A list of at least 5 queries (in natural language) you could write for your datasets. Some of these should require complex SQL (aggregations, subqueries, joins of 3 or more tables, etc.)

**Milestone 1 Submission Instructions:** One member should upload the proposal as PDF to Gradescope and add the other three members of the group to the submission as teammates.

Based on your project proposal, we will assign each group a TA who will serve as your project mentor for the remainder of the semester. Your project mentor will email your group with feedback on your proposal and add every member of your group to a private GitHub repository before the next milestone.

## Milestone 2 – Project Outline

Give a detailed functionality description and schema design, and initial set up.

In this phase, you will set up your version control environment, explain your project idea in more detail, and assign responsibilities to each group member.

First, make sure your project mentor has accepted the invitation into your GitHub repository. During development, you will use the repository to share code and perform version control. Your project mentor will also use the repository during grading to review the quality of your code and ensure you haven't plagiarized any of it.

Next, write your project outline. The outline should contain the following:

1. Motivation for the idea/description of the problem the application solves
2. List of features you will definitely implement in the application
3. List of features you might implement in the application, given enough time
4. List of pages the application will have and a 1-2 sentence description of each page. We expect that the functionality of each page will be meaningfully different from the functionality of the other pages. (At least 3 pages)
5. Relational schema as an ER diagram
6. SQL DDL for creating the database
7. Explanation of how you will clean and pre-process the data. [This tutorial](#) demonstrates how to do simple pre-processing in Python.
8. List of technologies you will use. You must use some kind of SQL database. We recommend using MySQL specifically because you will use MySQL in HW2, and we will provide guidance for setting up a MySQL database.
9. Description of what each group member will be responsible for

**Milestone 2 Submission Instructions:** One group member should then upload the project outline to Gradescope as a PDF and add all other group members to the submission as teammates

**IMPORTANT NOTE ABOUT AWS:** For the next milestone, you will populate and host your database on AWS. AWS database hosting isn't free. Each group will receive an email from AWS Academy, which grants \$100 in usage credits. You can follow the link in the email to register into Canvas and refer to the "AWS Academy Learner Lab Guide Student" pdf provided for further steps. This \$100 credit should be enough to complete the project in the semester. But if you exceed this amount, **you will be responsible for overages.**

You should not publicly share your id and password or put them somewhere that they can be compromised (Be careful about putting up the details in Github as well). We had an incident a few years ago where hackers used a group's AWS Keys to deploy several EC2 instances, which generated a charge of more than \$100 or some databases were compromised due to leak or password.

## Milestone 3 – Database Population

Pre-process your data, perform entity resolution, and populate the database.

Now that you have a schema, clean your data, perform entity resolution as needed, create a database and populate it with your data.

**Next, create a file containing the following:**

1. **Queries** -- A list of 5-10 SQL queries for your database that your application may execute. Ensure that 2-3 of your queries are complex (e.g. involve subqueries, joins, aggregations, etc.)  
For complex queries please ensure that you use them in your web app eventually - so make features which require complicated query processing.
2. **Descriptions** -- A 1-2 sentence description of what each query is supposed to do
3. **Credentials** -- Instructions and guest credentials for accessing the database

**Milestone 3 Submission Instructions:** One group member should upload the file of SQL queries to Gradescope and add all other group members to the submission as teammates

**Now, begin building your application!** Here are a few recommendations for getting started:

1. **Use the code from HW2 as a starting point.** Feel free to copy and paste it into your project directory. This doesn't count as plagiarism as long as you mention that you used it in your final report.
2. **Before building anything, make sure you can retrieve data from your database and display it on a webpage.** This requires setting up a server, connecting it to your database, and creating one route that queries your database then sends the data to the web client. You shouldn't need to write much code for any of these steps, but if you're not familiar with web development, putting all the pieces together can be challenging. Once you've done it once, it's easy to replicate and modify for your different features. Reach out to your project mentor early if you struggle with this step.



## Milestone 4 – API Specification and Project Mentor Check-In

### a. API Specification Document

Make a simple document containing the following for each route your app executes:

- Description of the functionality of the route
- Request path
- Request parameters, types, and descriptions
- Query Parameters
- Response parameters, types, and descriptions

You get to decide what request parameters are important based on what your query needs to execute properly. You can access the route parameters in `req.params` (if the param is a path param) or in `req.query` (if the param is a query param).

An effective way to define an API spec will include, for each route, the following:

1. The request path (ex. `/getUser`) along with the method (i.e. GET, POST, DELETE, PUT) and a description of what the route is doing (ex. retrieves a user by username).
2. The request params, including the name (ex. `username`), param type (path or query), data type (i.e. integer, string, etc.), required / optional indicator, and description (ex. `username` corresponding to a user)
3. The response params, including the name, data type, and description.

### b. Demo basic functionality

Schedule a 15-minute Zoom conference call with your project mentor and teammates to demonstrate that you have implemented a portion of the necessary features you listed in the project outline. Only one student needs to attend the call, though it is helpful if more than one attends. You don't need to prepare a formal presentation for this meeting. Just plan to walk your project mentor through the features you've implemented. (No document necessary for the demo)

**Your mentor will give you feedback on the features you've built and indicate whether they think your team is on track to finish a polished application.** If there are technical challenges your team has been unable to resolve, this is a good opportunity to ask your mentor for advice.

## Milestone 5 – Final Report, Code and Demo Video

Finish the application, record a demo of it, and write your report

### a. Final Report

Write a polished final report containing the following information:

1. **Title** -- Name of your website/application (Be creative!)
2. **Introduction** -- Explanation of project goals/target problem, description of application functionality, list of group members, motivation behind the project.
3. **Architecture** -- List of technologies, description of system architecture/application
4. **Data** -- For each dataset, a link to the source, a description of the data, relevant summary statistics, and an explanation of how you use the data
5. **Database** -- Explanation of data ingestion procedure and entity resolution efforts, ER diagram, number of instances in each table, normal form and justification.
6. **Web App description** - Brief pages and functionality overview. A 2-4 line description of pages implemented and how they are different.
7. **API Specification** - A list of routes and their associated descriptions, request parameters, and response parameters (from MS4).
8. **Queries** -- Examples of at least 5 queries in your application and explanations of how they're used. Please report the queries you think are most complex.
9. **Performance evaluation** -- recorded timings before and after optimization, descriptions of events being timed, and explanations of why caching, indexing, and other optimizations improved timings
10. **Technical challenges** -- List of technical challenges and how you overcame them

In most cases, the final report turns out to be about 6-10 pages. Feel free to copy text from your earlier submissions where appropriate, if you think it's sufficiently polished and clear.

### b. Code

Create a zip file containing all the code for the application. You should **NOT** include the node/Python packages your application depends on. You **SHOULD** include:

1. Application code
2. All code used for cleaning, wrangling, and ingesting the data
3. A list of dependencies  
(Essentially, the external libraries that you use and we need to NPM install in order to run your website. You could check the package.json file in the client folder of hw2)
4. Instructions for building it locally

### c. Demo

Use screen capture software such as Zoom to record a 2-4 minute video demo of your application. The instructors will only watch the first 4 minutes of your video, so don't make it longer. Your video should include demonstrations of all your application's pages and features. One or more of your group members should narrate it and describe the queries and optimizations associated with each feature. Describe how much each optimization affected performance.

Record timings associated with your application's features before and after performing optimizations, such as caching or indexing. For example, record how long it takes to execute queries, respond to client requests, and execute any other common tasks.

**Milestone 5 Submission Instructions:** Please submit each deliverable via Gradescope and add all team members to your group.

For the Demo video - please upload it onto the Penn google mail drive and ensure to make the link accessible to University of Pennsylvania - upload the link to Gradescope.

### Final Live Demo: (Last Week of Class)

Demo the final version of your application to your project mentor.

Your project mentor will schedule a 15 minute demo on Zoom with your full group. One of your group members will run the application locally and share their screen with the group. At least one of your group members will describe and demonstrate the application's main features. Then, your project mentor may ask questions about database or application design and query optimization. Your mentor may also ask to see how the website handles certain inputs. This video conference call will be recorded and saved.

# GRADING

## 1. Criteria

The project will be graded out of 100 points. The seven following criteria will be used to evaluate your work. The percentage next to each criteria indicates the portion of the final score attributed to that criteria.

### 1. **Technical Quality - 30%**

Does the application solve a non-trivial problem or answer non-trivial questions? Is the application robust and functional? Are some of the queries complex? Is the database designed well? Is the schema normalized? Was entity resolution and data cleaning performed correctly? Does the application perform interesting joins between data from different sources?

### 2. **Scope - 25%**

Does the application implement a sufficient number of features?  
Does the application have multiple pages? Does each page interact with a database? Are the datasets large?

### 3. **Presentation - 20%**

Is the final report written using clear and concise language and correct grammar? Is the final report well-organized and easy to read? Does the final report make use of charts and tables when appropriate? During the video demo, is the narrator easier to hear and understand? Is the information presented accurate?

### 4. **Optimization - 15%** - Attempted to optimize the queries using the concepts taught in class, and explained the results achieved.

Did the developers attempt to optimize their more complex queries using relevant optimization techniques (e.g. improving the query, caching, indexing)? Does the final report provide reasonable, correct explanations for why each optimization failed or succeeded? Does the final report include relevant and accurate pre and post optimization timings?

### 5. **Look and feel - 10%**

Is the application visually appealing? Is the user interface easy to use? Does the application make use of data visualizations and images when appropriate and feasible?

## 6. Participation

We expect all members of the team to contribute to the project, and that all members of the team be given the opportunity to contribute to the project.

The project will be graded according to the criteria above, and typically each team member will receive that as their individual project grade. However, in egregious cases where a team member does not contribute enough to the project (as measured by evidence on code in Github and commits, meetings with the project TA, as well as feedback from the other team members) **we will deduct a percentage of the overall project grade** from the team member's individual project grade.

Did team members work cooperatively and support each other? As an individual, did you complete your fair share of the work? As an individual, did you fully complete the tasks assigned to you by the group in a timely manner? As an individual, did you communicate with your group members regarding your ability to meet internal deadlines, as necessary?

## 2. Awards

The instructors will nominate and vote on several projects to receive awards at the end of the semester. Each project that wins at least one award will receive an additional two points of extra credit. You may obtain a final score greater than 100 with this extra credit.

The following awards will be distributed:

1. Best all-around application
2. Most beautiful application
3. Most technically complex application

# APPENDIX – RESOURCES

## 1. Dataset Resources

We've compiled a list of datasets and dataset aggregators to help jumpstart your search for data. But don't feel compelled to choose something off this list. This is just the tip of the iceberg. There's a lot more data out there!

### Dataset Aggregators

1. [Wikidata](#)
2. [DBpedia](#)
3. [World Bank Open Data](#)
4. [Tableau datasets](#)
5. [sqlbelle's list of datasets and dataset aggregators](#)
6. [AwesomeData's list of public datasets](#)
7. [fivethirtyeight datasets](#)
8. [OpenDataPhilly](#)
9. [NYC OpenData](#)
10. [Zillow Datasets](#)
11. [AWS Datasets](#)
12. [Google Cloud Datasets](#)
13. [Reddit Datasets](#)
14. [Buzzfeed Datasets](#)
15. [Data.world](#)
16. [Kaggle](#)
17. [Federal Government Datasets](#)
18. [UCI ML Repository](#)
19. [Academic Torrents](#)

## 2. Example Projects

The list below gives a few examples of projects people have done in the past in this course. We hope these help inspire your group and give you a better idea of what we're looking for. You can use the datasets, but must develop your own original application over the datasets using your own code. Also see the example videos in Coursera.

### World Bank database Factbook

The data made available by [WorldBank](#) presents the most current and accurate global development data available, and it includes national, regional and global estimates. A group of students used this data to make an application that displayed development data for selected countries/regions in interactive visualizations.

### World Travel Guide

A group used travel datasets like [Google's](#) to build an interface that lets users see information about places they could travel to, based on their travel preferences and needs.

### Soccer Fantasy League

A group used this [soccer dataset](#) along with data scraped from ESPN to build a Soccer fantasy league. They built interfaces that let users draft their own teams and implemented an account system to save teams and expose different features to different kinds of users.

### Olympics App

A group used the Summer Olympics dataset augmented with data from the WorldBank to create an application that helped people learn more about each country, while viewing how well that country has performed at the Olympics historically.

### Bike Rental Routes

A group developed a web-app that allowed people to find routes between two points in NYC using a combination of bike, subway and walking. Platforms such as Google Maps assume that the user either owns their own bike or has none at all. This project used the Citi [Bike database](#) JSON and combined that with [subway stations database](#) of New York City and calculated the shortest distance using a combination of all three, allowing the user to pick up a Citi Bike and drop it off during the journey.

### 3. Useful Tools

We've compiled a list of software tools that you may find helpful for scraping, cleaning, and ingesting data or making your demo video. If you know of some convenient tool we didn't include, please let us know on Piazza!

#### Content Extractors

1. [Apache Tika](#) -- Extracts metadata and text from word documents, PDF, powerpoints, and many other types of files
2. [Jackson Project](#) -- Efficient, easy-to-use JSON parser for Java
3. [Trail](#) -- XML Parser for Java
4. [jsoup](#) -- HTML parser for Java
5. [Selenium for Python](#) -- A Python library for automating interactions with a web browser to perform web scraping or application testing
6. [Beautiful Soup](#) -- A Python library for extracting data from HTML and XML files

#### Data cleaning and exploration tools

1. [Pandas](#) -- A Python library for data analysis that provides DataFrame, a convenient data structure for storing and processing large amounts of data
2. [Dora](#) -- A Python library for cleaning data and performing exploratory analysis
3. [JupyterLab](#) -- Software that enables you to work with Jupyter notebooks -- documents that contain live runnable Python code and display code output inline. Useful for data exploration, cleaning, and pre-processing
4. [Google Colab](#) -- An online version of JupyterLab that allows Google-Drivestyle collaboration

#### Screen Recording Software

1. [ScreenRec](#) -- Free screen recording software for Mac, Windows, and Linux
2. [Open Broadcaster Software](#) -- Free screen recording software for Mac, Windows, and Linux
3. [Zoom](#) can record calls, including multiple participants and shared screens

#### Database Software

1. [MySQL Workbench](#) -- GUI for MySQL
2. Oracle SQL



# PLAGIARISM POLICY

You can refer to the web or any other resource for ideas, but you are **STRICTLY NOT ALLOWED** to use other people's code directly. If you would like to use some code or snippets, please consult your project mentor and obtain permission before you do so.

Please make sure that you cite the original author/source if you are approved to use it.

If you are caught violating this policy, you will receive a 0 for the final project and/or the class, and we will refer your case to the Office of Student Conduct (OSC), which may take further disciplinary action.